

Representation & training

C5a: or HRMC6: collapse target rescC7s: WBM trained scalar
to residual cap (succ. .625→.975) is also viable

Planner integration

Discrete: additive Sarsa vs weights, focalC7a: C8: additive C13a: fresh certifier 0/6 →
L1 vs ManhattanC7b: 1.0: -6-15% equal vs loose Euclidean shared queue 6/6

Formulation, information supervision

C8/C9b: future wC9a/C9c: LoRA plateauC11/C12: no depthC13b: bounded obs. +
no benefit (0/9) capacity; labels set integrity dose-response Bellman: -15.95%