

Representation & training

C5: scalar HRM collapses
to residual cap

C6: field target rescues HRM
(succ. .625→.975)

C7: well-trained scalar
is also viable

Planner integration

Discrete: additive residual
inert vs Manhattan

Same weights, focal rank
 $w=1.0$: -6-15% exp.

C7-C8: additive beats
focal vs loose Euclid

C13: fresh certifier 0/6 →
shared queue 6/6

Formulation, information, supervision

C8/C9b: future window
no benefit (0/9)

C9/C9h: LoRA plateau =
capacity; labels set regime

C11/C12: no depth or
hierarchy dose-response

C13-M: bounded obs. +
local Bellman: -15.95%